

```
!pip install keras==2.10.0
!pip install tensorflow==2.10.0
!pip install h5py==3.7.0

Requirement already satisfied: keras==2.10.0 in /usr/local/lib/python3.10/dist-packages (2.10.0)
Requirement already satisfied: tensorflow==2.10.0 in /usr/local/lib/python3.10/dist-packages (2.10.0)
Requirement already satisfied: absl-py==1.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (1.4.0)
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (1.6.3)
Requirement already satisfied: flatbuffers>=2.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (23.5.26)
Requirement already satisfied: gast<=0.4.0,>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (0.4.0)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (0.2.0)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (1.56.0)
Requirement already satisfied: h5py>=2.9.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (3.7.0)
Requirement already satisfied: keras<2.11,>=2.10.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (2.10.0)
Requirement already satisfied: keras-preprocessing>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (1.1.0)
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (16.0.6)
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (1.22.4)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (3.3.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (23.1)
Requirement already satisfied: protobuf<3.20,>=3.9.2 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (3.19.6)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (67.7.2)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (1.16.0)
Requirement already satisfied: tensorboard<2.11,>=2.10 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (2.10.1)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0)
Requirement already satisfied: tensorflow-estimator<2.11,>=2.10.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (2.3.0)
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (4.7.1)
Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow==2.10.0) (1.14.1)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from astunparse>=1.6.0->tensorflow==2.10.0)
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.11,>=2.10->tensorflow)
Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.11,>=2.10->tensorflow)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.11,>=2.10->tensorflow)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.11,>=2.10->tensorflow)
Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.11,>=2.10->tensorflow)
Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.11,>=2.10->tensorflow)
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.11,>=2.10->tensorflow)
Requirement already satisfied: cachetools<6.0,>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3->tensorflow)
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3->tensorflow)
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3->tensorflow)
Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from google-auth-oauthlib<0.5,>=0.4.1->tensorflow)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorboard)
Requirement already satisfied: charset-normalizer~>2.0.0 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorboard<2.11,>=2.10->tensorflow)
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.10/dist-packages (from werkzeug>=1.0.1->tensorboard<2.11,>=2.10->tensorflow)
Requirement already satisfied: pyasn1<0.6.0,>=0.4.6 in /usr/local/lib/python3.10/dist-packages (from pyasn1-modules>=0.2.1->google-auth)
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.10/dist-packages (from requests-oauthlib>=0.7.0->google-auth)
Requirement already satisfied: h5py>=3.7.0 in /usr/local/lib/python3.10/dist-packages (3.7.0)
Requirement already satisfied: numpy>=1.14.5 in /usr/local/lib/python3.10/dist-packages (from h5py>=3.7.0) (1.22.0)
```

```
!pip install efficientnet
```

```
Requirement already satisfied: efficientnet in /usr/local/lib/python3.10/dist-packages (1.1.1)
Requirement already satisfied: keras-applications<=1.0.8,>=1.0.7 in /usr/local/lib/python3.10/dist-packages (from efficientnet) (1.0.7)
Requirement already satisfied: scikit-image in /usr/local/lib/python3.10/dist-packages (from efficientnet) (0.19.3)
Requirement already satisfied: numpy>=1.9.1 in /usr/local/lib/python3.10/dist-packages (from keras-applications<=1.0.8,>=1.0.7->efficientnet)
Requirement already satisfied: h5py in /usr/local/lib/python3.10/dist-packages (from keras-applications<=1.0.8,>=1.0.7->efficientnet)
Requirement already satisfied: scipy>=1.4.1 in /usr/local/lib/python3.10/dist-packages (from scikit-image->efficientnet) (1.10.1)
Requirement already satisfied: networkx>=2.2 in /usr/local/lib/python3.10/dist-packages (from scikit-image->efficientnet) (3.1)
Requirement already satisfied: pillow!=7.1.0,!=7.1.1,!=8.3.0,>=6.1.0 in /usr/local/lib/python3.10/dist-packages (from scikit-image->efficientnet)
Requirement already satisfied: imageio>=2.4.1 in /usr/local/lib/python3.10/dist-packages (from scikit-image->efficientnet) (2.4.1)
Requirement already satisfied: tifffile>=2019.7.26 in /usr/local/lib/python3.10/dist-packages (from scikit-image->efficientnet) (2019.7.26)
Requirement already satisfied: PyWavelets>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-image->efficientnet) (1.1.1)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from scikit-image->efficientnet) (23.1)
```

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).
```

```
#!unzip /content/drive/MyDrive/car_damage_classification/Data.zip -d /content/drive/MyDrive/car_damage_classification
```

```
Archive: /content/drive/MyDrive/car_damage_classification/Data.zip
  creating: /content/drive/MyDrive/car_damage_classification/Data/
  creating: /content/drive/MyDrive/car_damage_classification/Data/training/
  creating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/
  inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0001.jpeg
  inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0002.jpeg
  inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0003.jpeg
  inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0004.jpeg
  inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0005.jpeg
  inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0006.jpeg
  inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0007.jpeg
```

```
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0008.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0009.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0010.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0011.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0012.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0013.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0014.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0015.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0016.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0017.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0018.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0019.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0020.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0021.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0022.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0023.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0024.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0025.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0026.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0027.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0028.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0029.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0030.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0031.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0032.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0033.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0034.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0035.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0036.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0037.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0038.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0039.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0040.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0041.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0042.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0043.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0044.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0045.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0046.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0047.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0048.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0049.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0050.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0051.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0052.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0053.jpeg
inflating: /content/drive/MyDrive/car_damage_classification/Data/training/damage/0054.jpeg
```

```
import numpy as np
import pickle
import os
from tqdm import tqdm
import cv2
import pandas as pd
import seaborn as sns
from os import listdir
import tensorflow as tf
from sklearn.preprocessing import LabelBinarizer
from keras.models import Sequential
from sklearn.utils import shuffle
from keras.layers.convolutional import Conv2D
from keras.layers.convolutional import MaxPooling2D
from keras.layers.core import Activation, Flatten, Dropout, Dense
from keras import backend as K
from keras.preprocessing.image import ImageDataGenerator
from keras.optimizers import Adam, SGD
from keras.preprocessing import image
from tensorflow.keras.preprocessing.image import img_to_array
from keras.utils import to_categorical
import imblearn
from imblearn.over_sampling import SMOTE
from sklearn.preprocessing import MultiLabelBinarizer
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
from tensorflow.keras.models import load_model
from keras.models import Model
from tensorflow.keras.layers import Input
from keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPool2D, BatchNormalization, AveragePooling2D, GlobalAveragePooling2D
from keras.callbacks import ModelCheckpoint, ReduceLROnPlateau
from keras import layers
from sklearn.metrics import classification_report, confusion_matrix
import efficientnet.keras as effnet

from tensorflow.keras import backend as K
K.clear_session()
```

```

image_size=128
default_image_size = tuple((128, 128))
labels = os.listdir('/content/drive/MyDrive/car_damage_classification/Data/training')
directory_root = '/content/drive/MyDrive/car_damage_classification/Data'
classes = len(labels)
BATCHZ_SIZE=32

def image_array(dir):
    try:
        img = cv2.imread(dir)
        if img is not None :
            img = cv2.resize(img, default_image_size)
            return img_to_array(img)
        else :
            return np.array([])
    except Exception as e:
        print(f"Error : {e}")
        return None

image_list, label_list = [], []

try:
    print("[INFO] Loading images ...")
    root_dir = listdir(directory_root)
    for directory in root_dir :
        # remove .DS_Store from list
        if directory == ".DS_Store" :
            root_dir.remove(directory)

    for classes_folder in root_dir :
        image_classes_folder_list = listdir(f"/content/drive/MyDrive/car_damage_classification/Data/training")

    for classes_folder in image_classes_folder_list :
        # remove .DS_Store from list
        if classes_folder == ".DS_Store" :
            image_classes_folder_list.remove(classes_folder)

    for i in listdir('/content/drive/MyDrive/car_damage_classification/Data'):
        for image_classes_folder in image_classes_folder_list:
            print(f"[INFO] Processing {image_classes_folder} ...")
            image_classes_image_list = listdir(f"{directory_root}/{i}/{image_classes_folder}/")

            for single_image_classes_image in image_classes_image_list :
                if single_image_classes_image == ".DS_Store" :
                    image_classes_image_list.remove(single_image_classes_image)

            for image in image_classes_image_list[:]:
                image_directory = f"{directory_root}/{i}/{image_classes_folder}/{image}"
                if image_directory.endswith(".jpg") == True or image_directory.endswith(".JPG") == True or image_directory.endswith(".p"):
                    image_list.append(image_array(image_directory))
                    label_list.append( image_classes_folder)
            print("[INFO] Image loading completed")
except Exception as e:
    print(f"Error : {e}")

[INFO] Loading images ...
[INFO] Processing damage ...
[INFO] Processing nodamage ...
[INFO] Image loading completed
[INFO] Processing damage ...
[INFO] Processing nodamage ...
[INFO] Image loading completed

X = image_list
Y = label_list

X = np.array(X)

X, Y = shuffle(X, Y)
print(X.shape)

(2300, 128, 128, 3)


#Data Visualization
import matplotlib.image as mpimg
plt.figure(figsize = (20, 40))
image_count = 1

```

```
BASE_URL = '/content/drive/MyDrive/car_damage_classification/Data/training/'
for directory in os.listdir(BASE_URL):
    if directory[0] != '.':
        for i, file in enumerate(os.listdir(BASE_URL + directory)):
            if i == 1:
                break
            else:
                fig = plt.subplot(13, 3, image_count)
                image_count += 1
                image = mpimg.imread(BASE_URL + directory + '/' + file)
                plt.imshow(image)
                plt.title(directory)
```



```
#count plot
plt.figure(figsize = (12, 7))
plt.xticks(rotation=90)
sns.countplot(x=Y)
```



```
class_labels = LabelBinarizer()
Y = class_labels.fit_transform(Y)
pickle.dump(class_labels,open('/content/drive/MyDrive/car_damage_classification/label_transform.pkl', 'wb'))
n_classes = len(class_labels.classes_)

cls = len(class_labels.classes_)
print(class_labels.classes_)

['damage' 'nodamage']

Y = to_categorical(Y)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.1)

X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.1)
```

## CNN

```
#CNN
depth=3
model = Sequential()
inputShape = (image_size, image_size, depth)
chanDim = -1
if K.image_data_format() == "channels_first":
    inputShape = (depth, image_size, image_size)
    chanDim = 1
model.add(Conv2D(32, (3, 3), padding="same", input_shape=inputShape))
model.add(Activation("relu"))
model.add(BatchNormalization(axis=chanDim))
model.add(MaxPooling2D(pool_size=(3, 3)))
model.add(Dropout(0.25))
model.add(Conv2D(64, (3, 3), padding="same"))
model.add(Activation("relu"))
model.add(BatchNormalization(axis=chanDim))
model.add(Conv2D(64, (3, 3), padding="same"))
model.add(Activation("relu"))
model.add(BatchNormalization(axis=chanDim))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Conv2D(128, (3, 3), padding="same"))
model.add(Activation("relu"))
model.add(BatchNormalization(axis=chanDim))
model.add(Conv2D(128, (3, 3), padding="same"))
model.add(Activation("relu"))
model.add(BatchNormalization(axis=chanDim))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(1024))
model.add(Activation("relu"))
model.add(BatchNormalization())
model.add(Dropout(0.25))
model.add(Dense(cls))
model.add(Activation("softmax"))
model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d_5 (Conv2D)	(None, 128, 128, 32)	896
activation_7 (Activation)	(None, 128, 128, 32)	0
batch_normalization_6 (BatchNormalization)	(None, 128, 128, 32)	128
max_pooling2d_3 (MaxPooling2D)	(None, 42, 42, 32)	0
dropout_4 (Dropout)	(None, 42, 42, 32)	0
conv2d_6 (Conv2D)	(None, 42, 42, 64)	18496
activation_8 (Activation)	(None, 42, 42, 64)	0
batch_normalization_7 (BatchNormalization)	(None, 42, 42, 64)	256
conv2d_7 (Conv2D)	(None, 42, 42, 64)	36928
activation_9 (Activation)	(None, 42, 42, 64)	0
batch_normalization_8 (BatchNormalization)	(None, 42, 42, 64)	256
max_pooling2d_4 (MaxPooling2D)	(None, 21, 21, 64)	0
dropout_5 (Dropout)	(None, 21, 21, 64)	0
conv2d_8 (Conv2D)	(None, 21, 21, 128)	73856
activation_10 (Activation)	(None, 21, 21, 128)	0

```
batch_normalization_9 (BatchNormalization)      512
conv2d_9 (Conv2D)                            (None, 21, 21, 128)    147584
activation_11 (Activation)                   (None, 21, 21, 128)    0
batch_normalization_10 (BatchNormalization)    512
max_pooling2d_5 (MaxPooling2D)                (None, 10, 10, 128)    0
dropout_6 (Dropout)                          (None, 10, 10, 128)    0
flatten_1 (Flatten)                          (None, 12800)          0
dense_2 (Dense)                            (None, 1024)           13108224
```


```
model.compile(loss="categorical_crossentropy", optimizer= "adam", metrics=["accuracy"])
```

```
history = model.fit(X_train, y_train,batch_size=BATCH_SIZE,validation_data=(X_val, y_val),epochs=5, verbose=1)
```

```
Epoch 1/5
59/59 [=====] - 12s 41ms/step - loss: 0.8067 - accuracy: 0.7118 - val_loss: 0.6968 - val_accuracy: 0.6377
Epoch 2/5
59/59 [=====] - 2s 30ms/step - loss: 0.4380 - accuracy: 0.7923 - val_loss: 0.4913 - val_accuracy: 0.8019
Epoch 3/5
59/59 [=====] - 2s 29ms/step - loss: 0.4198 - accuracy: 0.8196 - val_loss: 0.6908 - val_accuracy: 0.7585
Epoch 4/5
59/59 [=====] - 2s 30ms/step - loss: 0.3732 - accuracy: 0.8465 - val_loss: 0.4925 - val_accuracy: 0.7923
Epoch 5/5
59/59 [=====] - 2s 31ms/step - loss: 0.3408 - accuracy: 0.8588 - val_loss: 0.6317 - val_accuracy: 0.7198
```

```
#accuracy and loss plot
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```

```
#loss plot
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```




```

pred = model.predict(X_test)
pred = np.argmax(pred, axis=1)
y_test_new = np.argmax(y_test, axis=1)

8/8 [=====] - 0s 30ms/step
model loss

cmat = confusion_matrix(y_test_new, pred)
plt.figure(figsize=(5,5))
sns.heatmap(cmat, annot = True, cbar = False, cmap='Blues', fmt="d");

```



```

print(classification_report(y_test_new,pred))

      precision    recall  f1-score   support

          0       0.68      0.91      0.78     112
          1       0.88      0.59      0.71     118

    accuracy                           0.75     230
   macro avg       0.78      0.75      0.74     230
weighted avg       0.78      0.75      0.74     230

```

```

from sklearn.metrics import precision_recall_fscore_support
res = []
for l in range(cls):
    prec,recall,_,_ = precision_recall_fscore_support(y_test_new==l,
                                                       pred==l,
                                                       pos_label=True,average=None)
    res.append([l,recall[0],recall[1]])

pd.DataFrame(res,columns = ['class','sensitivity','specificity'])

```

	class	sensitivity	specificity
0	0	0.593220	0.910714
1	1	0.910714	0.593220

## VGG16

```

from keras.applications import vgg16
def build_model(backbone):
    model = Sequential()
    model.add(backbone)
    model.add(layers.GlobalAveragePooling2D())
    model.add(layers.Dropout(0.5))
    model.add(layers.BatchNormalization())
    model.add(layers.Dense(1, activation='softmax'))

    model.compile(
        loss='binary_crossentropy',
        optimizer="adam",
        metrics=['accuracy']
    )

    return model

vgg_model = vgg16.VGG16(weights="imagenet",
                        include_top=False,
                        input_shape=(image_size,image_size,3))

base_model = vgg16.VGG16
trainable_layers = 4

base_model = base_model(weights="imagenet", include_top=False, input_shape=(image_size,image_size,3))

for layer in base_model.layers[:-trainable_layers]:
    layer.trainable = False

model = build_model(base_model)
model.summary()

Model: "sequential_3"

```

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 4, 4, 512)	14714688
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 512)	0
dropout_9 (Dropout)	(None, 512)	0
batch_normalization_13 (BatchNormalization)	(None, 512)	2048
dense_5 (Dense)	(None, 2)	1026

```

=====
Total params: 14,717,762
Trainable params: 7,081,474
Non-trainable params: 7,636,288

```

---

```

model.compile(loss='categorical_crossentropy', optimizer="adam", metrics=['accuracy'])

history = model.fit(X_train, y_train,batch_size=BATCHZ_SIZE,validation_data=(X_val, y_val),epochs=5, verbose=1)

Epoch 1/5
59/59 [=====] - 10s 120ms/step - loss: 0.4084 - accuracy: 0.8223 - val_loss: 0.4207 - val_accuracy: 0.8696
Epoch 2/5
59/59 [=====] - 3s 59ms/step - loss: 0.2838 - accuracy: 0.8808 - val_loss: 0.8029 - val_accuracy: 0.7681
Epoch 3/5
59/59 [=====] - 3s 59ms/step - loss: 0.2037 - accuracy: 0.9238 - val_loss: 3.5014 - val_accuracy: 0.6232
Epoch 4/5
59/59 [=====] - 3s 59ms/step - loss: 0.1882 - accuracy: 0.9238 - val_loss: 0.3209 - val_accuracy: 0.8889
Epoch 5/5
59/59 [=====] - 4s 61ms/step - loss: 0.1166 - accuracy: 0.9581 - val_loss: 0.6004 - val_accuracy: 0.7923

```



```

#accuracy and loss plot
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train','val'], loc='upper left')

```

```
plt.show()
```


```
#loss plot
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```



```
pred = model.predict(X_test)
pred = np.argmax(pred, axis=1)
y_test_new = np.argmax(y_test, axis=1)
```

```
8/8 [=====] - 1s 153ms/step
```

```
cmat = confusion_matrix(y_test_new, pred)
plt.figure(figsize=(5,5))
sns.heatmap(cmat, annot = True, cbar = False, cmap='Blues', fmt="d");
```



```
print(classification_report(y_test_new,pred))

      precision    recall  f1-score   support

          0       1.00     0.67     0.80     112
          1       0.76     1.00     0.86     118

   accuracy                           0.84    230
  macro avg       0.88     0.83     0.83    230
weighted avg       0.88     0.84     0.83    230
```

```
from sklearn.metrics import precision_recall_fscore_support
res = []
for l in range(cls):
    prec,recall,_,_ = precision_recall_fscore_support(y_test_new==l,
                                                       pred==l,
                                                       pos_label=True,average=None)
    res.append([l,recall[0],recall[1]])

pd.DataFrame(res,columns = ['class','sensitivity','specificity'])
```

	class	sensitivity	specificity
<b>0</b>	0	1.000000	0.669643
<b>1</b>	1	0.669643	1.000000

## Resnet50

```
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Flatten, GlobalAveragePooling2D, BatchNormalization, Dropout
from tensorflow.keras.applications import ResNet50
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, ReduceLROnPlateau, CSVLogger

resnet = ResNet50(weights= None, include_top=False, input_shape= X_train.shape[1:])

x = resnet.output
x = GlobalAveragePooling2D()(x)
x = Dropout(0.5)(x)
predictions = Dense(cls, activation= 'softmax')(x)
model = Model(inputs = resnet.input, outputs = predictions)
```

```
model.summary()
```

```
Model: "model"
```

Layer (type)	Output Shape	Param #	Connected to
<hr/>			
input_5 (InputLayer)	[(None, 128, 128, 3 0 )]	0	[]
conv1_pad (ZeroPadding2D)	(None, 134, 134, 3) 0	0	['input_5[0][0]']
conv1_conv (Conv2D)	(None, 64, 64, 64) 9472	9472	['conv1_pad[0][0]']
conv1_bn (BatchNormalization)	(None, 64, 64, 64) 256	256	['conv1_conv[0][0]']
conv1_relu (Activation)	(None, 64, 64, 64) 0	0	['conv1_bn[0][0]']
pool1_pad (ZeroPadding2D)	(None, 66, 66, 64) 0	0	['conv1_relu[0][0]']
pool1_pool (MaxPooling2D)	(None, 32, 32, 64) 0	0	['pool1_pad[0][0]']
conv2_block1_1_conv (Conv2D)	(None, 32, 32, 64) 4160	4160	['pool1_pool[0][0]']
conv2_block1_1_bn (BatchNormal	(None, 32, 32, 64) 256	256	['conv2_block1_1_conv[0][0]']

```

    ization)

conv2_block1_1_relu (Activation) (None, 32, 32, 64) 0      ['conv2_block1_1_bn[0][0]']

conv2_block1_2_conv (Conv2D) (None, 32, 32, 64) 36928     ['conv2_block1_1_relu[0][0]']

conv2_block1_2_bn (BatchNormal) (None, 32, 32, 64) 256     ['conv2_block1_2_conv[0][0]']

conv2_block1_2_relu (Activation) (None, 32, 32, 64) 0      ['conv2_block1_2_bn[0][0]']

conv2_block1_0_conv (Conv2D) (None, 32, 32, 256) 16640     ['pool1_pool[0][0]']

conv2_block1_3_conv (Conv2D) (None, 32, 32, 256) 16640     ['conv2_block1_2_relu[0][0]']

conv2_block1_0_bn (BatchNormal) (None, 32, 32, 256) 1024     ['conv2_block1_0_conv[0][0]']

conv2_block1_3_bn (BatchNormal) (None, 32, 32, 256) 1024     ['conv2_block1_3_conv[0][0]']

conv2_block1_add (Add) (None, 32, 32, 256) 0      ['conv2_block1_0_bn[0][0]', 'conv2_block1_3_bn[0][0]']

conv2_block1_out (Activation) (None, 32, 32, 256) 0      ['conv2_block1_add[0][0]']

conv2_block2_1_conv (Conv2D) (None, 32, 32, 64) 16448     ['conv2_block1_out[0][0]']

conv2_block2_1_bn (BatchNormal) (None, 32, 32, 64) 256     ['conv2_block2_1_conv[0][0]']

conv2_block2_1_relu (Activation) (None, 32, 32, 64) 0      ['conv2_block2_1_bn[0][0]']

```

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model.fit(X_train, y_train, batch_size=BATCH_SIZE, validation_data=(X_val, y_val), epochs=5, verbose=1)
```


```

Epoch 1/5
59/59 [=====] - 18s 164ms/step - loss: 1.1320 - accuracy: 0.6286 - val_loss: 10.8030 - val_accuracy: 0.570
Epoch 2/5
59/59 [=====] - 7s 125ms/step - loss: 0.7408 - accuracy: 0.6656 - val_loss: 0.7248 - val_accuracy: 0.5169
Epoch 3/5
59/59 [=====] - 7s 127ms/step - loss: 0.6772 - accuracy: 0.7198 - val_loss: 0.7562 - val_accuracy: 0.5556
Epoch 4/5
59/59 [=====] - 7s 126ms/step - loss: 0.6627 - accuracy: 0.7327 - val_loss: 358.4745 - val_accuracy: 0.507
Epoch 5/5
59/59 [=====] - 8s 128ms/step - loss: 0.5959 - accuracy: 0.7654 - val_loss: 0.8404 - val_accuracy: 0.5314

```

```
#accuracy and loss plot
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```

```
#loss plot
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```



model loss




```

pred = model.predict(X_test)
pred = np.argmax(pred, axis=1)
y_test_new = np.argmax(y_test, axis=1)

8/8 [=====] - 2s 93ms/step
      8          /          V          I
cmat = confusion_matrix(y_test_new, pred)
plt.figure(figsize=(5,5))
sns.heatmap(cmat, annot = True, cbar = False, cmap='Blues', fmt="d");

```



```
print(classification_report(y_test_new,pred))
```

	precision	recall	f1-score	support
0	0.89	0.07	0.13	112
1	0.53	0.99	0.69	118
accuracy			0.54	230
macro avg	0.71	0.53	0.41	230
weighted avg	0.70	0.54	0.42	230

```

from sklearn.metrics import precision_recall_fscore_support
res = []
for l in range(cls):
    prec,recall,_,_ = precision_recall_fscore_support(y_test_new==l,
                                                       pred==l,
                                                       pos_label=True,average=None)
    res.append([l,recall[0],recall[1]])

pd.DataFrame(res,columns = ['class','sensitivity','specificity'])

```

	class	sensitivity	specificity
0	0	0.991525	0.071429
1	1	0.071429	0.991525

## EfficientNetB2

```

base_model = effnet.EfficientNetB2(weights='imagenet', include_top=False, input_shape=(image_size, image_size, 3))
model = base_model.output
model = GlobalAveragePooling2D()(model)
model = Dense(cls, activation='softmax')(model)
model = Model(inputs = base_model.input, outputs=model)

Downloading data from https://github.com/Callidior/keras-applications/releases/download/efficientnet/efficientnet-b2\_weights\_tf\_dim\_31936256/31936256 [=====] - 0s 0us/step

```



```

model.compile(loss='categorical_crossentropy', optimizer=Adam(), metrics=['accuracy'])

model.summary()

```

Layer (type)	Output Shape	Param #	Connected to
input_6 (InputLayer)	[(None, 128, 128, 3 0 )]	0	[]
stem_conv (Conv2D)	(None, 64, 64, 32)	864	['input_6[0][0]']
stem_bn (BatchNormalization)	(None, 64, 64, 32)	128	['stem_conv[0][0]']
stem_activation (Activation)	(None, 64, 64, 32)	0	['stem_bn[0][0]']
block1a_dwconv (DepthwiseConv2 D)	(None, 64, 64, 32)	288	['stem_activation[0][0]']
block1a_bn (BatchNormalization )	(None, 64, 64, 32)	128	['block1a_dwconv[0][0]']
block1a_activation (Activation )	(None, 64, 64, 32)	0	['block1a_bn[0][0]']
block1a_se_squeeze (GlobalAver agePooling2D)	(None, 32)	0	['block1a_activation[0][0]']
block1a_se_reshape (Reshape)	(None, 1, 1, 32)	0	['block1a_se_squeeze[0][0]']
block1a_se_reduce (Conv2D)	(None, 1, 1, 8)	264	['block1a_se_reshape[0][0]']
block1a_se_expand (Conv2D)	(None, 1, 1, 32)	288	['block1a_se_reduce[0][0]']
block1a_se_excite (Multiply)	(None, 64, 64, 32)	0	['block1a_activation[0][0]', 'block1a_se_expand[0][0]']
block1a_project_conv (Conv2D)	(None, 64, 64, 16)	512	['block1a_se_excite[0][0]']
block1a_project_bn (BatchNorma lization)	(None, 64, 64, 16)	64	['block1a_project_conv[0][0]']
block1b_dwconv (DepthwiseConv2 D)	(None, 64, 64, 16)	144	['block1a_project_bn[0][0]']
block1b_bn (BatchNormalization )	(None, 64, 64, 16)	64	['block1b_dwconv[0][0]']
block1b_activation (Activation )	(None, 64, 64, 16)	0	['block1b_bn[0][0]']
block1b_se_squeeze (GlobalAver agePooling2D)	(None, 16)	0	['block1b_activation[0][0]']
block1b_se_reshape (Reshape)	(None, 1, 1, 16)	0	['block1b_se_squeeze[0][0]']

```
block1b_se_reduce (Conv2D)      (None, 1, 1, 4)      68      ['block1b_se_reshape[0][0]']
block1b_se_expand (Conv2D)      (None, 1, 1, 16)     80      ['block1b_se_reduce[0][0]']
block1b_se_excite (Multiply)   (None, 64, 64, 16)    0      ['block1b_activation[0][0]',


history = model.fit(X_train, y_train,batch_size=BATCHZ_SIZE,validation_data=(X_val, y_val),epochs=5, verbose=1)

Epoch 1/5
59/59 [=====] - 31s 229ms/step - loss: 0.3415 - accuracy: 0.8696 - val_loss: 1.2843 - val_accuracy: 0.5845
Epoch 2/5
59/59 [=====] - 10s 167ms/step - loss: 0.1674 - accuracy: 0.9442 - val_loss: 0.7193 - val_accuracy: 0.7923
Epoch 3/5
59/59 [=====] - 9s 154ms/step - loss: 0.1057 - accuracy: 0.9608 - val_loss: 0.4156 - val_accuracy: 0.9034
Epoch 4/5
59/59 [=====] - 9s 160ms/step - loss: 0.0855 - accuracy: 0.9758 - val_loss: 0.4839 - val_accuracy: 0.8937
Epoch 5/5
59/59 [=====] - 11s 180ms/step - loss: 0.0432 - accuracy: 0.9871 - val_loss: 0.4087 - val_accuracy: 0.9034

#model.save('/content/drive/MyDrive/car_damage_classification/Models/EfficientnetB2.h5')

#accuracy and loss plot
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()

#loss plot
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```




```

pred = model.predict(X_test)
pred = np.argmax(pred, axis=1)
y_test_new = np.argmax(y_test, axis=1)

8/8 [=====] - 3s 107ms/step
|          |
cmat = confusion_matrix(y_test_new, pred)
plt.figure(figsize=(5,5))
sns.heatmap(cmat, annot = True, cbar = False, cmap='Blues', fmt="d");

```



```

print(classification_report(y_test_new,pred))

      precision    recall   f1-score   support
0       0.96     0.88     0.92     112
1       0.90     0.97     0.93     118

accuracy                           0.93    230
macro avg       0.93     0.93     0.93    230
weighted avg    0.93     0.93     0.93    230

```

```

from sklearn.metrics import precision_recall_fscore_support
res = []
for l in range(cls):
    prec,recall,_,_ = precision_recall_fscore_support(y_test_new==l,
                                                       pred==l,
                                                       pos_label=True,average=None)
    res.append([l,recall[0],recall[1]])

pd.DataFrame(res,columns = ['class','sensitivity','specificity'])

```

	class	sensitivity	specificity
0	0	0.966102	0.883929
1	1	0.883929	0.966102

## Testing

```

class_labels = pd.read_pickle(r'/content/drive/MyDrive/car_damage_classification/label_transform.pkl')
label = class_labels.classes_.tolist()
label

['damage', 'nodamage']

```

```
image = cv2.imread('/content/drive/MyDrive/car_damage_classification/Data/validation/damage/0001.jpeg')
image = img_to_array(image)
image = cv2.resize(image, (image_size, image_size))
image = np.array([image])

image.shape
(1, 128, 128, 3)

model = load_model('/content/drive/MyDrive/car_damage_classification/Models/EfficientnetB2.h5')
prediction=model.predict(image)
print(prediction)
print(max(prediction[0]))

1/1 [=====] - 3s 3s/step
[[9.9910200e-01 8.9798676e-04]]
0.999102

pred_= prediction[0]
pred=[]
for ele in pred_:
    pred.append(ele)

maxi_ele = max(pred)
idx = pred.index(maxi_ele)
final_class=label
class_name= final_class[idx]
print("Predicted Class Is : " + str(class_name))

Predicted Class Is : damage
```